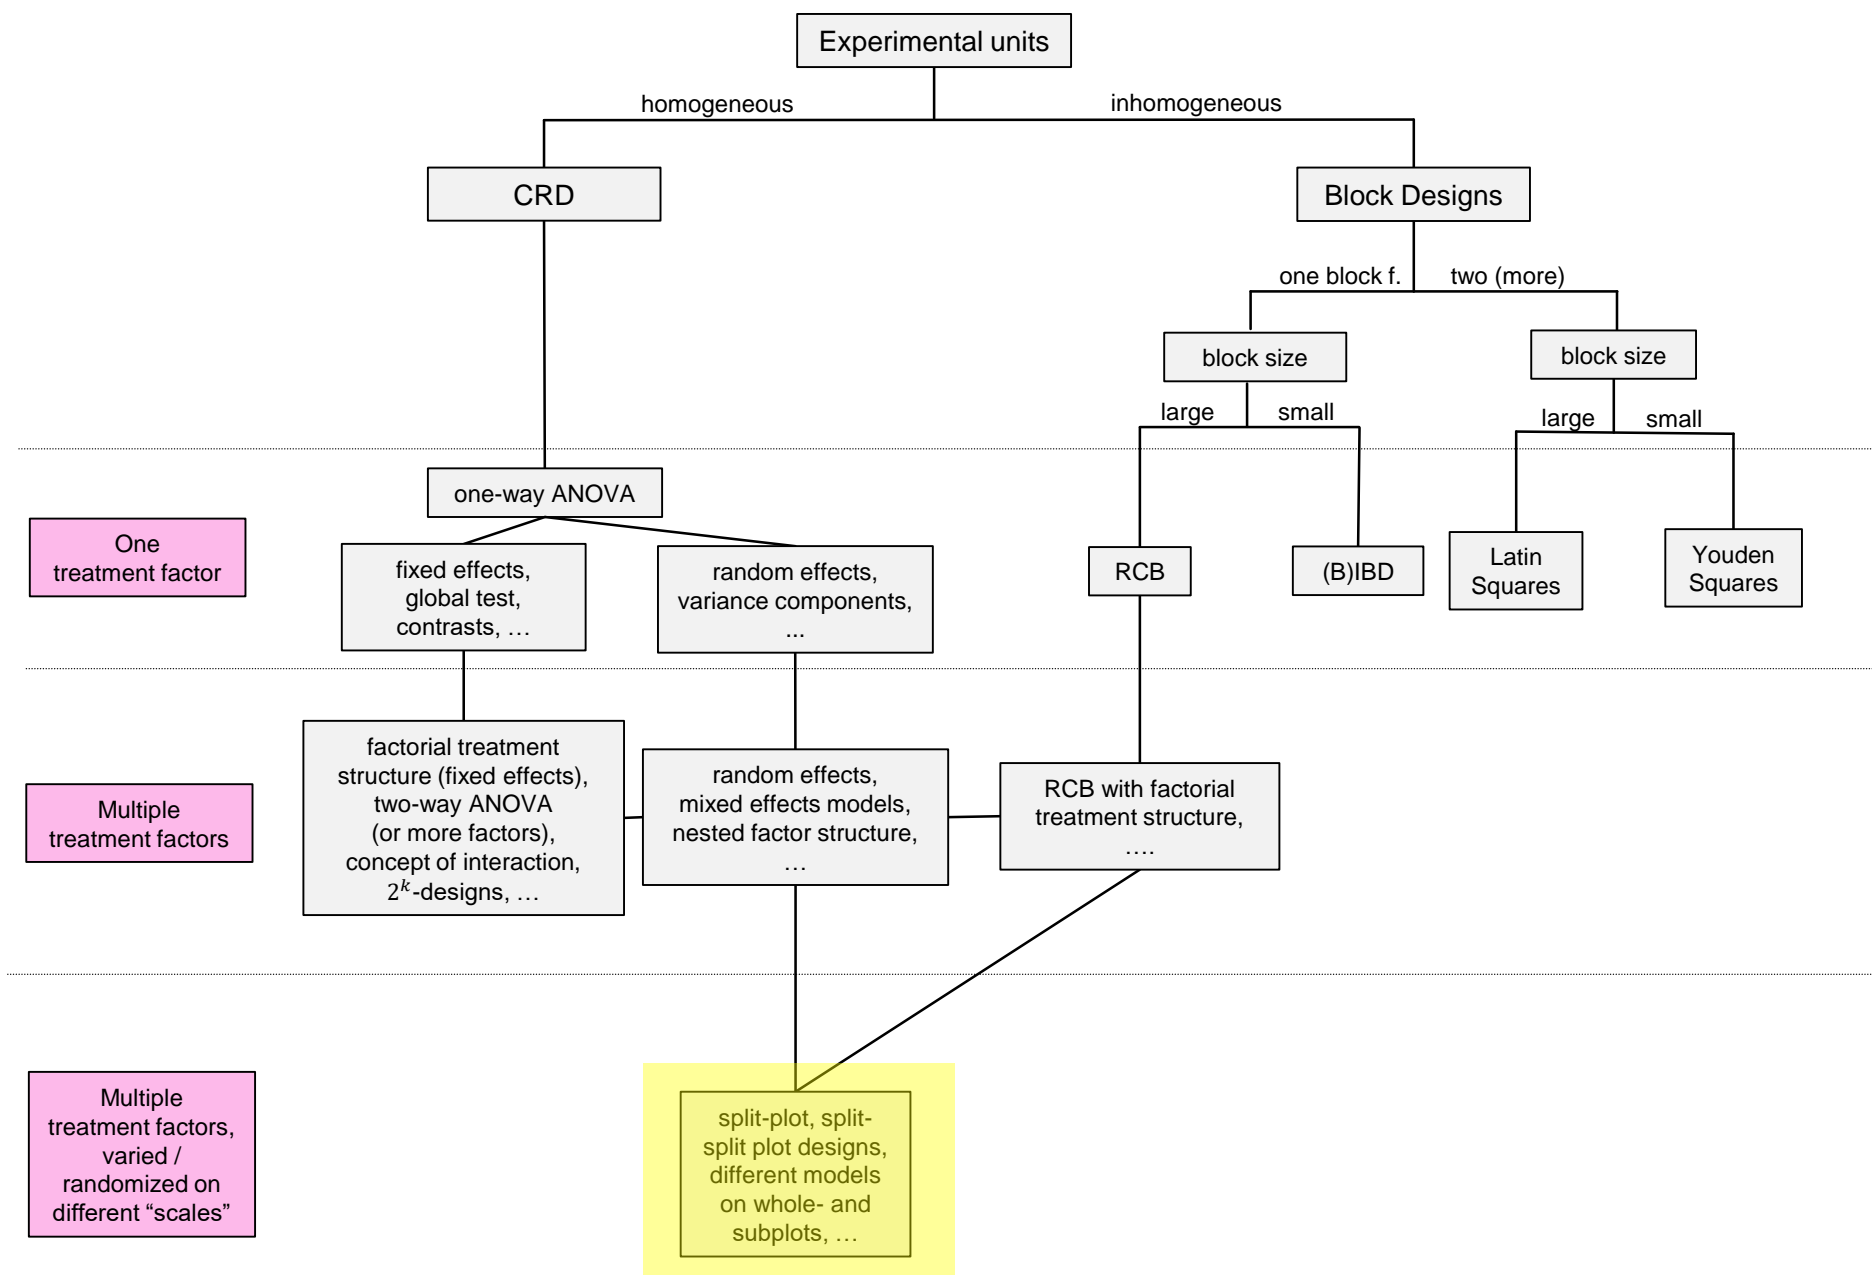


10



Split-Plot Designs



What is a Split-Plot Design? (Oehlert, 2010, Chapter 16.1)

- A split-plot design is a **special case** of a design with **factorial treatment structure**.
- It is used when some factors are **harder** (or more expensive) to **vary** than others.
- Basically, the standard split-plot design consists of two experiments with **different experimental units** of different “sizes”.
- E.g., in agronomic field trials certain factors require “large” experimental units, whereas other factors can be easily applied to “smaller” plots of land.
- Let us have a look at an example ...

Example I: Irrigation and Corn Variety (Oehlert, 2010)



- Consider the following factorial problem:
 - 3 different **irrigation levels**
 - 4 different **corn varieties**
 - Response: **Biomass**
 - Available resources: **6 plots of land**
- By definition, we **cannot** vary the irrigation level on a too small scale.
- We are “forced” to use “large” experimental units for the irrigation level factor.
- Assume that we can use one specific irrigation level on each of the 6 plots.

Example I: Irrigation and Corn Variety

- We randomly assign each irrigation level to 2 of the plots, the so called **whole plots** or **main plots**.
- In each of the plots, we randomly assign the 4 different corn varieties to the so called **split plots** (or sub plots).

4	4	1	3	2	2
2	1	4	1	4	1
3	3	2	2	1	4
1	2	3	4	3	3

- **Two independent randomizations** are being performed!
- We also call irrigation level the **whole-plot factor** and corn variety the **split-plot factor**.

Example I: Irrigation and Corn Variety

- **Whole plots** (plots of land) are the experimental units for the **whole-plot factor** (irrigation level).
- **Split plots** (subplots of land) are the experimental units for the **split-plot factor** (corn variety).
- In the split-plot “world”, whole plots act as **blocks**.
- Basically, we are performing **two different experiments in one**:
 - each experiment has its own **randomization**,
 - each experiment has its own idea of **experimental unit**.

Example I: Irrigation and Corn Variety

- How can we model this type of data?
- We use a **mixed model** formulation with **two different errors**:

$$Y_{ijk} = \mu + \alpha_i + \eta_{k(i)} + \beta_j + (\alpha\beta)_{ij} + \varepsilon_{k(ij)}$$

Diagram illustrating the mixed model equation and its components:

- Y_{ijk} is labeled as **biomass**.
- μ is the overall mean.
- α_i is the **fixed effect of irrigation**.
- $\eta_{k(i)}$ is the **whole-plot error**, distributed as $N(0, \sigma_{\eta}^2)$.
- β_j is the **fixed effect of corn variety**.
- $(\alpha\beta)_{ij}$ is the **(fixed) interaction between irrigation and corn variety**.
- $\varepsilon_{k(ij)}$ is the **split-plot error**, distributed as $N(0, \sigma^2)$.

- This means: Observations in the same whole plot share the same whole-plot error $\eta_{k(i)}$ and are therefore **not independent**.
- In R, this model can be easily fitted using `lmer` with a **random effect** (better terminology here: **error**) of the form `(1 | whole.plot)`.

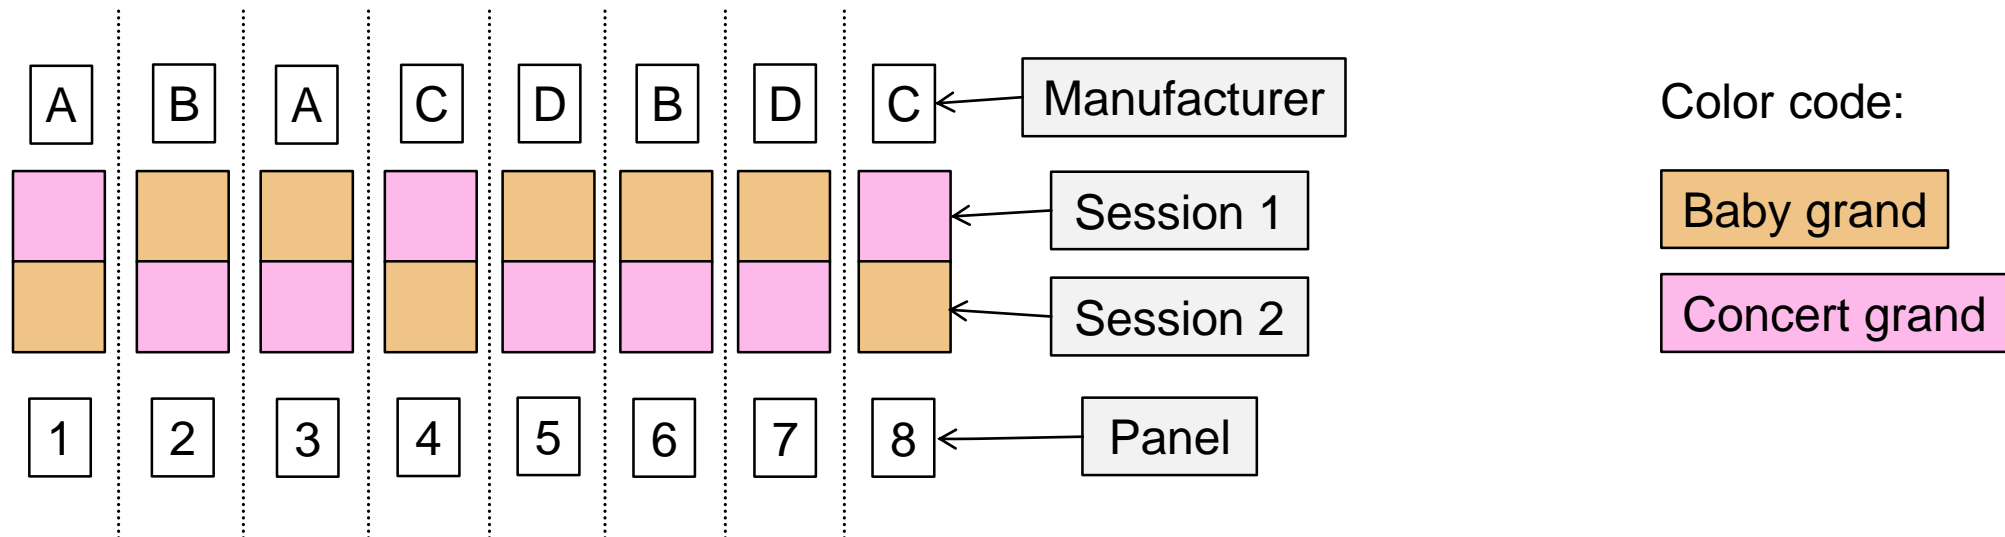
Example II: Pianos (Oehlert, 2010)



- **2 piano types** (“baby grand” & “concert grand”) from each of **4 manufacturers** (A, B, C, D).
- 40 music students are divided at random into **8 groups** (“panels”) of 5 students each.
- 2 panels are assigned at random to each manufacturer (= 2 panels per manufacturer).
- Each panel goes to the concert hall and hears (blindfolded) the sound of **both pianos** (in random order).
- **Response:** Average rating of the 5 students in the panel (hence, a student is “only” a measurement unit here).

Example II: Pianos

- The **whole plots** are the 8 **panels**.
- The **whole-plot factor** is the **manufacturer**.
- The **split plots** are the 2 **session time-slots**.
- The **split-plot factor** is the **piano type** (baby vs. concert grand).



Example II: Pianos

- The model is the same:

$$Y_{ijk} = \mu + \alpha_i + \eta_{k(i)} + \beta_j + (\alpha\beta)_{ij} + \varepsilon_{k(ij)}$$

Diagram illustrating the components of the mixed-effects model for piano ratings:

- μ : average rating of panel
- α_i : fixed effect of manufacturer
- $\eta_{k(i)}$: whole-plot error (distributed as $N(0, \sigma_\eta^2)$)
- β_j : fixed effect of piano type
- $(\alpha\beta)_{ij}$: (fixed) interaction between manufacturer and piano type
- $\varepsilon_{k(ij)}$: split-plot error (distributed as $N(0, \sigma^2)$)

- Again: This means that observations in the same whole-plot share the same whole-plot error $\eta_{k(i)}$ and are therefore **not independent**.

Example III: Oats



- Dataset `oats` from R-package `MASS`.
- As stated in the help file:

*The yield of oats from a **split-plot field trial** using 3 varieties and 4 levels of manurial treatment. The experiment was laid out in 6 blocks of 3 main plots, each split into 4 sub-plots. The varieties were applied to the main plots and the manurial treatments to the sub-plots.*
- Overview of data:
 - 6 different **blocks** (B)
 - 3 different **varieties** (V)
 - 4 different **nitrogen treatments** (N)
 - **Response** (Y): Yields (in $\frac{1}{4}$ lbs per sub-plot, each of area $\frac{1}{80}$ acre).
- Let us first have a **graphical overview** of the experimental design.

Example III: Oats

4	4	1
2	1	3
3	3	2
1	2	4

I

2	1	1
1	2	4
3	4	2
4	3	3

II

3	3	2
2	2	3
1	4	4
4	1	1

III

1	1	3
2	3	2
4	2	1
3	4	4

IV

3	4	3
2	1	4
4	2	1
1	3	2

V

2	3	1
1	4	4
4	2	2
3	1	3

VI

Example III: Oats

- This is a more complicated design as before as we have an additional **block factor**.
- A **whole plot** is given by a plot of land in a block.
- The **whole-plot factor** is variety.
- A **block design (RCB)** was used at the whole-plot level.
- A **split plot** is given by a **subplot of land**.
- The **split-plot factor** is given by **nitrogen treatment**.

Example III: Oats

- We have a **RCBD** for the whole-plot factor.
- An **experimental unit** on the whole-plot level is given by the **combination of block and variety**.
- We therefore use the model

$$Y_{ijk} = \mu + \alpha_i + \gamma_k + \eta_{ik} + \beta_j + (\alpha\beta)_{ij} + \varepsilon_{ijk}$$

The diagram illustrates the components of the mixed-effects model for the Oats experiment. It shows the following mappings:

- yield** (response variable) points to Y_{ijk} .
- fixed effect of variety** points to α_i .
- fixed effect of block** points to γ_k .
- whole-plot error** points to η_{ik} , which is associated with the random effect $N(0, \sigma_\eta^2)$.
- fixed effect of nitrogen treatment** points to β_j .
- (fixed) interaction between variety and nitrogen treatment** points to $(\alpha\beta)_{ij}$.
- split-plot error** points to ε_{ijk} , which is associated with the random effect $N(0, \sigma^2)$.

Example III: Oats

- In R we use the `lmer` function with an extra random effect (better terminology here: **error**) per **combination** of **block** and **variety**.

- We get the following output:

```
> fit.lme <- lmer(Y ~ B + V * N + (1 | B:V), data = oats)
> anova(fit.lme)
```

Analysis of Variance Table of type III with Satterthwaite approximation for degrees of freedom

	Sum Sq	Mean Sq	NumDF	DenDF	F.value	Pr(>F)	
B	4675.0	935.0	5	10	5.280	0.01244	*
V	526.1	263.0	2	10	1.485	0.27239	
N	20020.5	6673.5	3	45	37.686	2.458e-12	***
V:N	321.8	53.6	6	45	0.303	0.93220	

- Observe that the test for variety uses 2 and 10 degrees of freedom, respectively.
- Why? Let us have a closer look at the potential ANOVA table on the whole-plot level.

Example III: Oats

- On the whole-plot level we have the following ANOVA table:

<i>Source</i>	<i>df</i>
Block	5
Variety	2
Error (whole-plot)	10 (= 17 – 7)
Total	17 (= 18 – 1)

- Think of “averaging away” the nitrogen factor, hence we have **one observation per combination of block and variety**.
- Technically speaking, variety is tested against the **interaction** of block and variety.

Example III: Oats

- This also reveals a problem: We don't have too many error df's left to test the main effect of the whole-plot factor (only 10).
- In contrast, we test everything involving the **split-plot** factor against the **residual error**, which has 45 df's.

- Remember:

```
> fit.lme <- lmer(Y ~ B + V * N + (1 | B:V), data = oats)
> anova(fit.lme)
```

Analysis of Variance Table of type III with Satterthwaite approximation for degrees of freedom

	Sum Sq	Mean Sq	NumDF	DenDF	F.value	Pr(>F)	
B	4675.0	935.0	5	10	5.280	0.01244	*
V	526.1	263.0	2	10	1.485	0.27239	
N	20020.5	6673.5	3	45	37.686	2.458e-12	***
V:N	321.8	53.6	6	45	0.303	0.93220	

- Hence, the main effect of the whole-plot factor is estimated **less precisely** and the test is **less powerful** (compared to the split-plot level).
→ “Prize” for “laziness” on the whole-plot level (only a few observations).

General Situation

- Split-plot designs can also arise in (much) **more complicated** forms.
- There can be **more than one whole-plot factor**. E.g., think of a two-way factorial on the whole-plot level.
- In addition, there can be **more than one factor on the split-plot level**.
- To get the correct model we “only” have to **follow “the path of randomization”**.
- For each “level” (whole plot / split plot) of the experiment we have to introduce a corresponding random effect (better terminology here: **error**) which acts as the experimental error on that level.

General Situation

- This means:
 - Start on the whole-plot level and forget about the split plots.
 - Write down the corresponding model equation (incl. random effect / error).
 - Move on to the next level, expand equation with new terms (the upper level is now a block).
 - Etc.
- In R we just have to make sure that we tell `lmer` the correct random effects.
- In R it is sometimes useful to **define new variables** which identify the different experimental units on the different levels in order to specify the random effect structure.

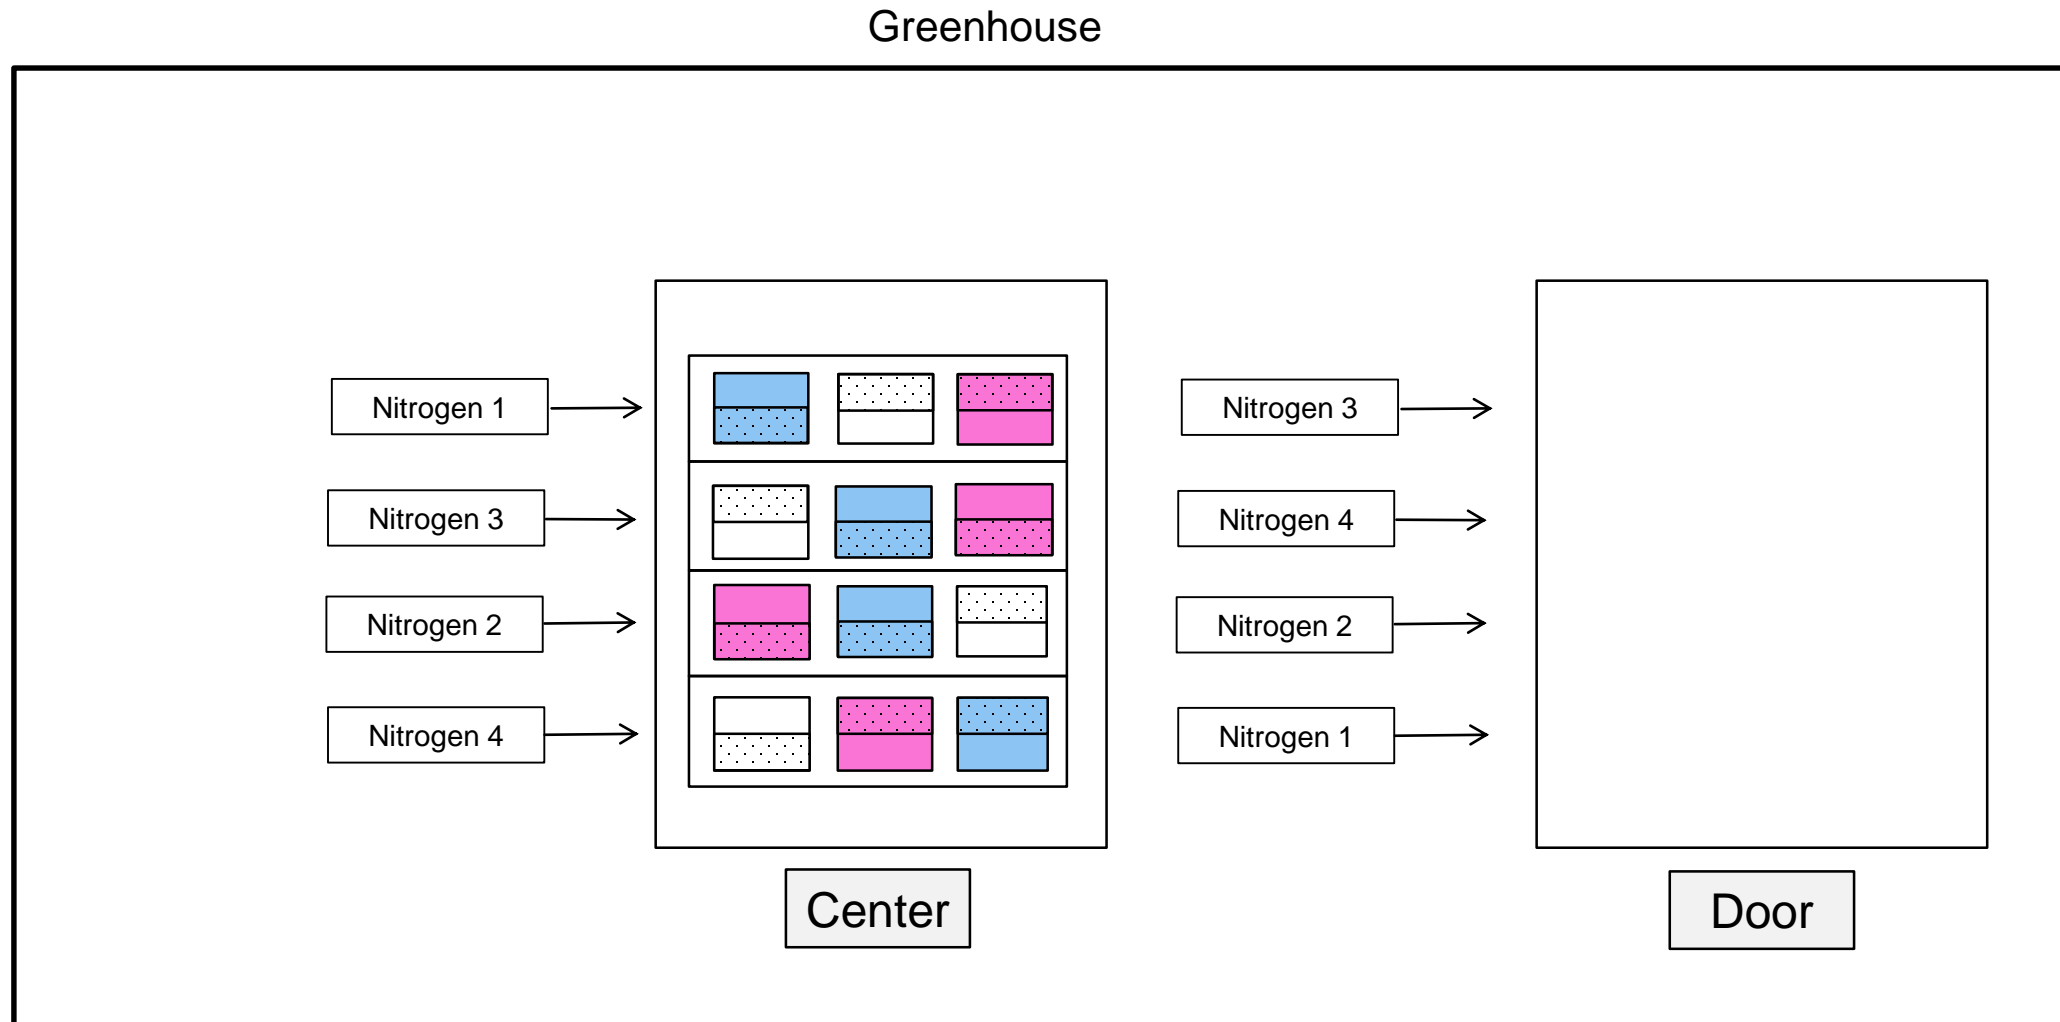
Example IV: Weed Biomass in Wetlands (Oehlert, 2010, Ex. 16.7)



- The experiment studied the effect of
 - **nitrogen** (4 levels of nitrogen)
 - **weed** (3 levels)
 - **clipping treatments** (2 levels: clipping / no clipping)on **plant growth** in wetlands.
- The experiment was performed as follows:
 - 8 **trays**, whereof each holds 3 artificial wetlands (rectangular wire baskets)
 - 4 of the trays were placed on **a table near the door** of the greenhouse,
 - 4 of the trays on **a table in the center** of the greenhouse.
 - On **each table**, we randomly assign one of the trays to each of the 4 **nitrogen treatments**.
 - Within **each tray**, we randomly assign the 3 **weed treatments**.
 - In addition, each wetland is split in half. One half is chosen at random and will be clipped, the other half is not clipped.
 - After 8 weeks: Measure fraction of biomass that is nonweed.

Example IV: Weed Biomass in Wetlands

Experimental layout



Example IV: Weed Biomass in Wetlands

- Let us follow the path of randomization:
 - Position in the greenhouse is a **block factor** (with levels center / door).
 - **Trays** are **whole plots**, and **nitrogen level** is the **whole-plot factor**.
 - **Wetlands** are **split plots** and **weed treatment** is the **split-plot factor**.
 - **Wetland halves** are so called **split-split plots** and **clipping** is the **split-split-plot factor**.
- Hence, we have a so-called **split-split plot design**.
- Let us now try to fit a model to this data set in R.

Example IV: Weed Biomass in Wetlands

- We use the following model:

```
> fit <- lmer(pct.nonweed.biomass ~ table + (1 | tray) + weed * nitrogen * clipping + (1 | wetland),  
+ data = wetland)
```

```
> anova(fit)
```

Type III Analysis of Variance Table with Satterthwaite's method

	Sum Sq	Mean Sq	NumDF	DenDF	F value	Pr(>F)	
table	0.16	0.16	1	3	0.1538	0.72113	
weed	1186.82	593.41	2	8	555.4531	2.613e-09	***
nitrogen	36.73	12.24	3	3	11.4610	0.03765	*
clipping	125.45	125.45	1	12	117.4290	1.494e-07	***
weed:nitrogen	157.57	26.26	6	8	24.5814	9.665e-05	***
weed:clipping	0.25	0.12	2	12	0.1149	0.89246	
nitrogen:clipping	0.74	0.25	3	12	0.2293	0.87419	
weed:nitrogen:clipping	4.82	0.80	6	12	0.7514	0.62033	

- All **main effects** and the **nitrogen × weed interaction** are significant.

Example IV: Weed Biomass in Wetlands

- We are here performing three experiments in one.
- On the **whole-plot level** we have the “experiment”:

<i>Source</i>	<i>df</i>
Table (block)	1
Nitrogen	3
Error (per tray)	3 (= 7 – 4)
Total	7 (= 8 – 1)

- On the **split-plot level** we have the “experiment”:

<i>Source</i>	<i>df</i>
Block (=Tray)	7
Weed	2
Weed × Nitrogen	6
Error (per wetland)	8 (= 23 – 15)
Total	23 (= 24 – 1)

Example IV: Weed Biomass in Wetlands

- On the **split-split-plot level** we have the “experiment”:

<i>Source</i>	<i>df</i>
Block (= wetland)	23
Clipping	1
Weed × Clipping	2
Nitrogen × Clipping	3
Nitrogen × Weed × Clipping	6
Error (per wetland half)	12 (= 47 – 35)
Total	47 (= 48 – 1)

Summary

- Split-plot designs and more complicated versions thereof are useful if some factors are harder (more expensive, ...) to vary than others.
- The “prize” for “laziness”: Less precision (power) on whole-plot level.
- To identify the correct design we **have to know the randomization procedure**.
- The general situation can be very complex, but by following the different randomization levels / steps, setting up a model is **easy**.
- Mixed effects software like `lmer` automatically identify the correct denominator for tests if the random effects / errors are stated correctly.