

Evolutionary game theory

Niko Beerenwinkel




Outline


- Evolutionary games
- Evolutionary stability
- The replicator equation
- The Prisoner's Dilemma

Constant selection

- So far, we have assumed that fitness is a constant parameter (*constant selection*).
- Example: Type A can move faster than type B.





fitness of A = 1.1



fitness of B = 1

Frequency-dependent selection

- But the selective advantage might actually decrease with increasing density of type A (e.g., because faster movements are impossible due to blocked space).



Frequency-dependent selection between two types

- Consider two types A and B, with frequencies x_A and x_B , respectively.
- The vector $x(t) = (x_A(t), x_B(t))^T$ describes the population.
- Denote by $f_A(x(t))$ the fitness of A and by $f_B(x(t))$ the fitness of B. The average fitness is $\phi(x) = x_A f_A(x) + x_B f_B(x)$.
- The deterministic selection dynamics are given by

$$\dot{x}_A = x_A [f_A(x) - \phi(x)]$$

$$\dot{x}_B = x_B [f_B(x) - \phi(x)]$$

Equilibria


- With $x := x_A$ and $1 - x = x_B$, the system is equivalent to

$$\dot{x} = x(1 - x) [f_A(x) - f_B(x)]$$

- Equilibria: $x = 0$, $x = 1$, and $\{x \in (0, 1) \mid f_A(x) = f_B(x)\}$
- $x^* = 0$ is stable if $f_A(0) < f_B(0)$
- $x^* = 1$ is stable if $f_A(1) > f_B(1)$
- An interior equilibrium x^* is stable if

$$\frac{\partial f_A}{\partial x}(x^*) < \frac{\partial f_B}{\partial x}(x^*)$$

Stability of equilibria



x ... abundance of A
 $1 - x$... abundance of B

$f_A(x) - f_B(x)$
 fitness difference between A and B

selection dynamics

● ... stable equilibrium

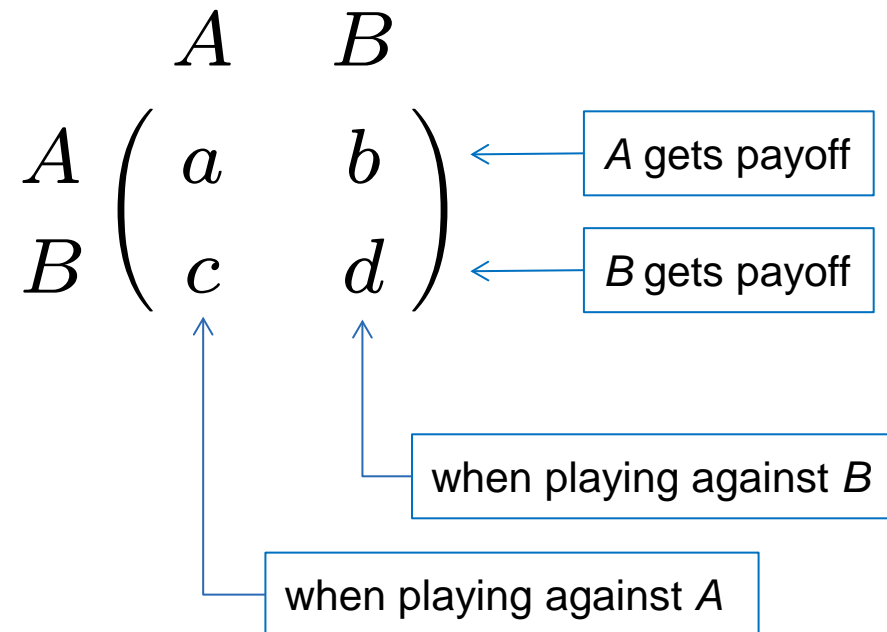
○ ... unstable equilibrium

The general idea of evolutionary game theory

- Evolutionary game theory is the study of *frequency-dependent selection*.
- We consider a population of players.
- Each player has a fixed strategy.
- Players interact (play the game) randomly.
- Success in the game is translated into reproductive success: Good strategies reproduce faster.

Evolutionary games with two players

- A two-player game is defined by the payoff matrix



Fitness = expected payoff

- We define fitness as the expected payoff,

$$f_A(x_A, x_B) = ax_A + bx_B$$

$$f_B(x_A, x_B) = cx_A + dx_B$$

- Setting $x := x_A$ we have $1 - x = x_B$ and we obtain the selection dynamics

$$\dot{x} = x(1 - x) [(a - b - c + d)x + b - d]$$

$$\begin{matrix} & A & B \\ \begin{matrix} A \\ B \end{matrix} & \begin{pmatrix} a & b \\ c & d \end{pmatrix} \end{matrix}$$

(i) A dominates B, if $a > c$ and $b > d$:



(ii) B dominates A, if $a < c$ and $b < d$:



(iii) A and B are bistable, if $a > c$ and $b < d$:



(iv) A and B coexist, if $a < c$ and $b > d$:



(v) A and B are neutral, if $a = c$ and $b = d$:



A is always the best strategy.
A is always fitter than B.

B is always the best strategy.
B is always fitter than B.

Playing the same strategy than
your opponent is the best strategy.
 $x^* = (d - b)/(a - b - c + d)$

Playing the opposite strategy than
your opponent is the best strategy.
 $x^* = (d - b)/(a - b - c + d)$

No matter what you do, the two
payoffs are always the same.

Nash equilibrium

- **Definition:** If two players play the same strategy and neither player can increase its payoff by changing strategy, then the strategy is at *Nash equilibrium*. Mathematically,
 - A is a strict Nash equilibrium, if $a > c$.
 - A is a Nash equilibrium, if $a \geq c$.
 - B is a strict Nash equilibrium, if $d > b$.
 - B is a Nash equilibrium, if $d \geq b$.
- $$\begin{array}{cc} & A & B \\ \begin{array}{c} A \\ B \end{array} & \left(\begin{array}{cc} a & b \\ c & d \end{array} \right) \end{array}$$
- A Nash equilibrium is the best reply to itself.

Examples

$$\begin{array}{cc} & A & B \\ A & \left(\begin{array}{cc} 3 & 0 \end{array} \right) \\ B & \left(\begin{array}{cc} 5 & 1 \end{array} \right) \end{array}$$

A is not Nash.

B is Nash.

B dominates A.

Note: Both playing B
has lower total payoff
than both playing A.

$$\begin{array}{cc} & A & B \\ A & \left(\begin{array}{cc} 3 & 1 \end{array} \right) \\ B & \left(\begin{array}{cc} 5 & 0 \end{array} \right) \end{array}$$

A is not Nash.

B is not Nash.

$$\begin{array}{cc} & A & B \\ A & \left(\begin{array}{cc} 5 & 0 \end{array} \right) \\ B & \left(\begin{array}{cc} 3 & 1 \end{array} \right) \end{array}$$

A is Nash.

B is Nash.

Evolutionary stable strategy (ESS)

- Suppose an infinitesimally small quantity ϵ of B invades an all-A population. Then selection will oppose invasion, if

$$f_A(1 - \epsilon) > f_B(\epsilon)$$

$$a(1 - \epsilon) + b\epsilon > c(1 - \epsilon) + d\epsilon$$

i.e., if in the limit as $\epsilon \rightarrow 0$

$$(a > c) \text{ or } (a = c \text{ and } b > d)$$

- **Definition:** A is ESS if either (i) $a > c$, or
(ii) $a = c$ and $b > d$

More than two strategies

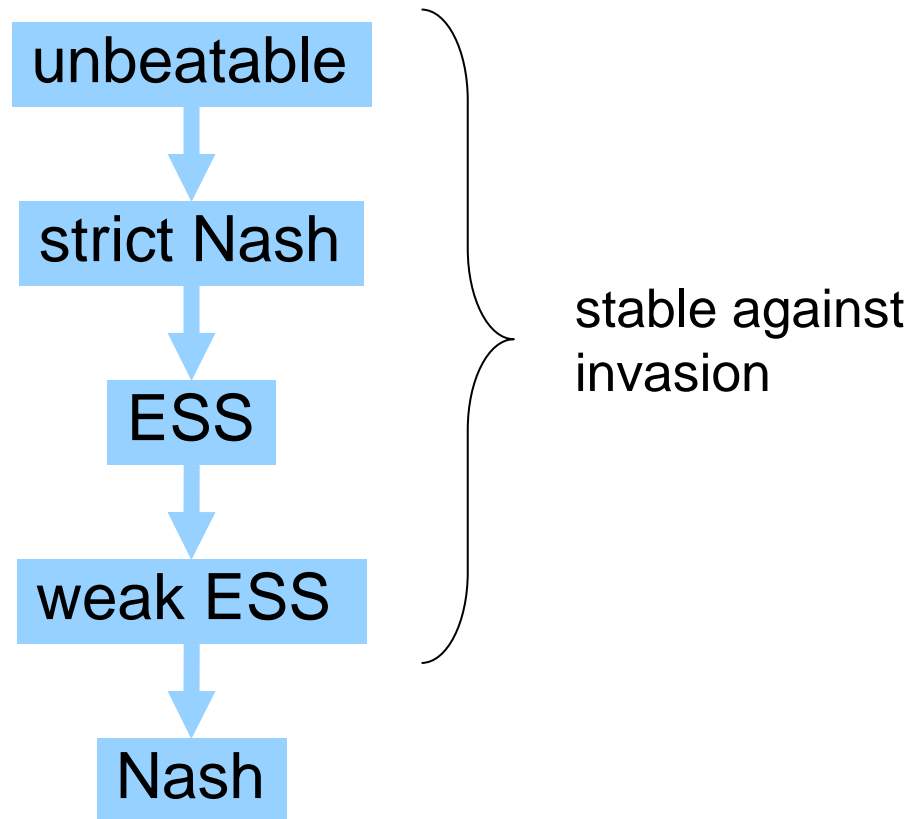
$$\begin{array}{c} S_1 \\ S_2 \\ \vdots \\ S_n \end{array} \begin{pmatrix} S_1 & S_2 & \dots & S_n \\ a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & \ddots & & \\ \vdots & & & \vdots \\ a_{n1} & & \dots & a_{nn} \end{pmatrix}$$

- The payoff for strategy S_i versus S_j is $E(S_i, S_j) = a_{ij}$.

Generalized definitions

- S_k is a *strict Nash* equilibrium, if $a_{kk} > a_{ik}$ for all $i \neq k$.
- S_k is a *Nash* equilibrium, if $a_{kk} \geq a_{ik}$ for all $i \neq k$.
- S_k is *ESS* if for all $i \neq k$, either
 - $a_{kk} > a_{ik}$, or
 - $a_{kk} = a_{ik}$ and $a_{ki} > a_{ii}$
- S_k is *weak ESS* if for all $i \neq k$, either
 - $a_{kk} > a_{ik}$, or
 - $a_{kk} = a_{ik}$ and $a_{ki} \geq a_{ii}$
- S_k is *unbeatable* if for all $i \neq k$, both
 - $a_{kk} > a_{ik}$, **and**
 - $a_{ki} > a_{ii}$

Hierarchy of concepts



Frequency-dependent fitness among n players

- Fitness = expected payoff:

$$f_i(x) = f_{S_i}(x) = \sum_{j=1}^n x_j a_{ij}$$

- Average population fitness:


$$\phi(x) = \sum_{i=1}^n x_i f_i(x)$$

The replicator equation


$$\dot{x}_i = x_i [f_i(x) - \phi(x)] \quad i = 1, \dots, n$$
$$x_1 + \dots + x_n = 1$$

- Fitness is frequency-dependent, usually in a linear fashion
- Generalizes constant selection dynamics
- The interior of the simplex S_n is invariant.
- Each face is invariant.
- Depending on the payoff matrix there can be fixed points in the interior and in every face of the simplex.

Three mating strategies among side-blotched lizards



- Orange-throated males
 - strongest
 - do not form strong pair bonds
 - fight blue-throated males for their females




- Blue-throated males
 - middle-sized
 - form strong pair bonds
 - can defend against yellow-throated males




- Yellow-throated males
 - smallest
 - color mimics females
 - can approach orange-throated males in disguise and mate with their females while engaged in fighting

Cyclic dominance



Rock-paper-scissors (RPC)



Payoff matrix for the RPC game

- We can re-scale any column of the payoff matrix.
- For the symmetric RPC game, we have

$$\begin{array}{c} R \\ S \\ P \end{array} \begin{array}{ccc} R & S & P \\ \left(\begin{array}{ccc} 0 & 1 & -1 \\ -1 & 0 & 1 \\ 1 & -1 & 0 \end{array} \right) \end{array}$$

- RPC is a zero-sum game with average fitness $\phi = 0$


Replicator dynamics of the RPC game

- The general RPC game is defined by the payoff matrix

$$A = \begin{pmatrix} 0 & -a_2 & b_3 \\ b_1 & 0 & -a_3 \\ -a_1 & b_2 & 0 \end{pmatrix}$$


- Case 1: $\det(A) > 0$
 - Unique interior equilibrium which is globally stable
 - Trajectories converge to this point as damped oscillations
- Case 2: $\det(A) < 0$
 - Unique interior equilibrium which is unstable
 - Trajectories converge to the boundary of the simplex in oscillations with increasing amplitude

$$\det(A) = 0$$




neutral oscillations

$$\det(A) > 0$$





damped oscillations
to a stable equilibrium

$$\det(A) < 0$$



increasing oscillations
to a heteroclinic cycle

What happened to the lizards?



More than three strategies

- For $n \geq 4$, the replicator equation allows for limit cycles and chaotic attractors.
- An interior equilibrium is given by the solution of the linear system $f_1 = \dots = f_n$ and $x_1 + \dots + x_n = 1$, which has one or zero non-degenerate solutions. There can be at most one isolated equilibrium in the interior.
- If there is no interior equilibrium, then all trajectories converge to the boundary of the simplex.

Conventional fights

- Many conflicts between animals do not escalate.
- In a *conventional fight*, threatening signals are exchanged that allow the contestants to assess each other's strength or determination, before one of them simply walks away.



Hawks and doves

- Let us analyze conventional fighting from the perspective of **individual selection** (rather than group selection).
- There are two basic strategies: Hawks (H) escalate fights, whereas doves (D) retreat.
- The benefit of winning a fight is b .
- The cost of injury is c .

$$\begin{array}{cc} & \begin{array}{cc} H & D \end{array} \\ \begin{array}{c} H \\ D \end{array} & \left(\begin{array}{cc} (b - c)/2 & b \\ 0 & b/2 \end{array} \right) \end{array}$$

Hawk and dove replicator dynamics

- Often, the cost of injury will be larger than the benefit of escalation ($b < c$).
- If $b < c$, then neither strategy is a Nash equilibrium.
- Hawks and doves can coexist. The equilibrium hawk frequency is

$$x_H^* = \frac{b}{c}$$

- If $b \ll c$, then there will be few hawks.

Mixed strategies

- Consider a strategy that plays “hawk” with probability p and “dove” with probability $1 - p$.
- The space of strategies is the interval $[0, 1]$, rather than the discrete space $\{H, D\}$ as before.
- The payoff for strategy p_1 versus strategy p_2 is

$$\begin{aligned} E(p_1, p_2) = & p_1 p_2 E(H, H) \\ & + p_1 (1 - p_2) E(H, D) \\ & + (1 - p_1) p_2 E(D, H) \\ & + (1 - p_1) (1 - p_2) E(D, D) \end{aligned}$$

Mixed hawk and dove strategy

- We find

$$E(p_1, p_2) = \frac{b}{2} \left(1 + p_1 - p_2 - \frac{c}{b} p_1 p_2 \right)$$

- The strategy $p^* = b/c$ is evolutionary stable:

$$E(p^*, p^*) = (b/2)[1 - (b/c)] = E(p, p^*)$$

$$E(p^*, p) = (b/2)[1 + (b/c) - 2p]$$

$$E(p, p) = (b/2)[1 - (c/b)p^2]$$

Thus p^* is Nash, not strict Nash, and ESS, because $E(p^*, p) > E(p, p)$ for all $p \neq p^*$.

There is always a Nash equilibrium

- For a given $n \times n$ payoff matrix A , the payoff for strategy $q \in S_n$ versus strategy $p \in S_n$ is

$$E(q, p) = qAp = \sum_{i=1}^n \sum_{j=1}^n a_{ij} q_i p_j$$

- There exists at least one strategy q^* such that

$$q^* A q^* \geq p A q^* \text{ for all } p$$

i.e., q^* is a Nash equilibrium.

The prisoner's dilemma

- Two people are suspected of having committed a joint crime. The state attorney offers separately the same deal to both: confess, become a witness, and avoid prison.
- Payoff = – years in prison

	remain silent	confess
remain silent	-1	-10
confess	0	-7

- Rational analysis of the game leads to “confess” no matter what the partner does. They will not cooperate.

The Prisoner's Dilemma captures the essential problem of cooperation

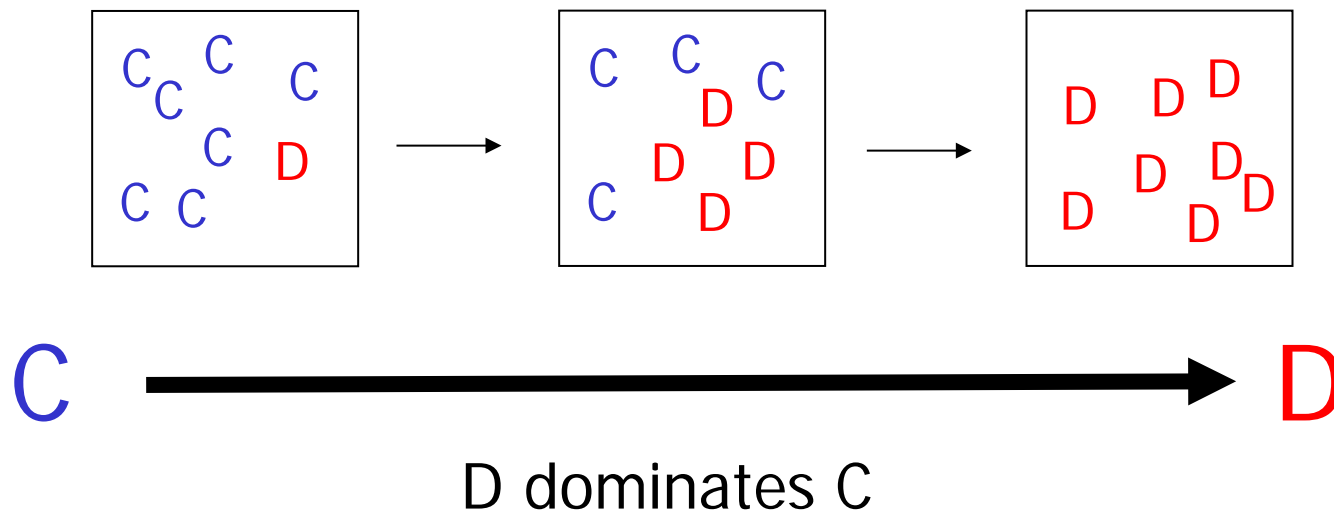
$$\begin{array}{cc} & C & D \\ \begin{array}{c} C \\ D \end{array} & \left(\begin{array}{cc} 3 & 0 \\ 5 & 1 \end{array} \right) \end{array}$$

- Rational players defect in order to maximize their payoff.
- Yet, cooperation is abundant in nature. For example,
 - among organelles in a cell
 - among cells maintaining a tissue
 - among tissues forming an organ
 - among organs in a body
 - among bacteria and animals
 -

Can cooperation evolve?

- Let us consider the replicator equation.
- $f_C(x) = 3x < f_D(x) = 5x + 1 - x = 4x + 1$
- Thus natural selection favors defection.

$$\begin{array}{c} C \\ D \end{array} \begin{array}{cc} C & D \\ \left(\begin{array}{cc} 3 & 0 \\ 5 & 1 \end{array} \right) \end{array}$$



The general Prisoner's Dilemma game

$$\begin{array}{cc} & C & D \\ \begin{array}{c} C \\ D \end{array} & \left(\begin{array}{cc} R & S \\ T & P \end{array} \right) \end{array}$$

$$T > R > P > S \quad \text{and} \quad R > \frac{T + P}{2}$$

- DC: Temptation to defect
- CC: Reward for mutual cooperation
- DD: Punishment for mutual defection
- CD: Sucker's payoff

Direct reciprocity

- The Prisoner's Dilemma game is repeated m times.
- Consider two strategies:
 - GRIM: Cooperate initially, then as long as opponent does not defect
 - ALLD: Always defect

	GRIM	ALLD
GRIM	mR	$S + (m - 1)P$
ALLD	$T + (m - 1)P$	mP

- If $m > (T - P)/(R - P)$, then GRIM is strict Nash, ALLD cannot invade. However, ALLD is also strict Nash.
- Hence, direct reciprocity can only stabilize cooperation.

Defecting in the last round

- If you know that you play m rounds, it is rational to defect in the last round. Use $(\)^*$ to denote this variant strategy.

	GRIM	GRIM [*]
GRIM	mR	$(m - 1)R + S$
GRIM [*]	$(m - 1)R + T$	$(m - 1)R + P$

- GRIM^{*} dominates GRIM, GRIM \rightarrow GRIM^{*}.
- But then:

$$\text{GRIM} \rightarrow \text{GRIM}^* \rightarrow \text{GRIM}^{**} \rightarrow \dots \rightarrow \text{GRIM}^{*(m)} = \text{ALLD}$$

Variable number of rounds

- Suppose that after each round there is a probability w that another round will be played.
- The probability of playing k rounds is

$$w^{k-1} (1 - w)$$

- Thus, the expected number of rounds is

$$\begin{aligned}\bar{m} &= (1 - w) + 2w(1 - w) + 3w^2(1 - w) + \dots \\ &= 1 - w + 2w - 2w^2 + 3w^2 - 3w^3 + \dots \\ &= 1 + w + w^2 + w^3 + \dots = \frac{1}{1 - w}\end{aligned}$$

Prisoner's Dilemma with a variable number of rounds

- The payoff for GRIM versus ALLD is

	GRIM	ALLD
GRIM	$\bar{m}R$	$S + (\bar{m} - 1)P$
ALLD	$T + (\bar{m} - 1)P$	$\bar{m}P$

- GRIM is evolutionary stable if $\bar{m} > (T - P)/(R - P)$.
- Defecting in the last round is no longer possible.

Tit-for-tat

- TFT strategy: start with cooperation, then do whatever your opponent has done in the previous round, i.e., answer C for C and D for D.
- TFT has won *Axelrod's tournaments*, an all-against-all competition between game strategies.
- Properties of TFT:
 - TFT is “nice”: it is never first to defect.
 - TFT is not greedy: it will never have higher payoff than a single opponent.
 - TFT is cooperative.

TFT versus ALLD

- I play TFT : C C C C C ...
You play TFT : C C C C C ...
- I play TFT : C D D D D ...
You play ALLD : D D D D D ...

	TFT	ALLD
TFT	$\bar{m}R$	$S + (\bar{m} - 1)P$
ALLD	$T + (\bar{m} - 1)P$	$\bar{m}P$

- TFT can resist invasion by ALLD if $\bar{m} > (T - P)/(R - P)$.
- TFT can resume cooperation, if the opponent cooperates.

TFT in the presence of noise

- Suppose errors occur in playing TFT.

TFT	:	C	C	C	[*] D	C	D	C	D	D	D	...
TFT	:	C	C	C	C	D	C	D	[*] D	D	D	...

- In the long run, the payoff is as low as choosing randomly between C and D,

$$E(\text{TFT}, \text{TFT}) = (T + R + P + S)/4 < R$$

because $R > (T + P)/2$ and $R > P$.


TFT versus ALLC

$$\begin{array}{cc}
 & \begin{array}{cc} \text{TFT} & \text{ALLC} \end{array} \\
 \begin{array}{c} \text{TFT} \\ \text{ALLC} \end{array} & \left(\begin{array}{cc} \bar{m}R & \bar{m}R \\ \bar{m}R & \bar{m}R \end{array} \right)
 \end{array}$$

- TFT is not evolutionary stable.
- In a finite population, random drift can lead from TFT to ALLC.

Reactive strategies

- A reactive strategy is a probabilistic strategy that takes into account the opponent's action in the previous round.
- The strategy $S(p, q)$ cooperates with
 - probability p if the opponent has cooperated in the previous move,
 - probability q if the opponent has defected in the previous move.
- Thus, p and q are conditional probabilities and each reactive strategy $S(p, q)$ corresponds to a point (p, q) in the unit square.



Reactive strategies for the repeated PD

- The repeated Prisoner's Dilemma between two reactive strategies $S_1(p_1, q_1)$ and $S_2(p_2, q_2)$ is a Markov chain with state space $\{CC, CD, DC, DD\}$, where, e.g., “CD” means “I play C and you play D”.
- The transition probability matrix M is

$$\begin{array}{c}
 \begin{array}{cccc}
 & CC & CD & DC & DD \\
 CC & \begin{pmatrix} p_1 p_2 & p_1(1-p_2) & (1-p_1)p_2 & (1-p_1)(1-p_2) \\
 CD & q_1 p_2 & q_1(1-p_2) & (1-q_1)p_2 & (1-q_1)(1-p_2) \\
 DC & p_1 q_2 & p_1(1-q_2) & (1-p_1)q_2 & (1-p_1)(1-q_2) \\
 DD & q_1 q_2 & q_1(1-q_2) & (1-q_1)q_2 & (1-q_1)(1-q_2)
 \end{pmatrix}
 \end{array}
 \end{array}$$

Stationary distribution of the Markov chain


- Let $x(t) = (x_{CC}(t), x_{CD}(t), x_{DC}(t), x_{DD}(t))$ be the probability distribution of the game after t rounds.
- We have $x(t + 1) = x(t) \cdot M$.
- The payoff at the stationary distribution is

$$E(S_1, S_2) = Rs_1s_2 + Ss_1(1 - s_2) + \\ + T(1 - s_1)s_2 + P(1 - s_1)(1 - s_2)$$

where $s_1(p_1, p_2, q_1, q_2)$ and $s_2(p_1, p_2, q_1, q_2)$ are the probabilities that players 1 and 2, respectively, cooperate in the stationary distribution.

Generous Tit-for-tat (GTFT)

- For the Prisoner's Dilemma with $R=3$, $T=5$, $S=0$, $P=1$, GTFT is defined by $p=1$ and $q=1/3$, i.e., $\text{GTFT} = S(1, 1/3)$.



Generous Tit-for-tat is more forgiving

- GTFT is more forgiving than TFT: it will cooperate one out of three times if the opponent has defected.
- $E(\text{GTFT}, \text{GTFT})$ is close to R , because GTFT can correct mistakes:

GTFT : C C C ^{*}D C D C C C C ...
GTFT : C C C C D C C C C C ...

- $E(\text{GTFT}, \text{ALLD}) < E(\text{TFT}, \text{ALLD})$, but in simulations of evolving strategies we often observe

ALLD \rightarrow TFT \rightarrow GTFT

GTFT for the general Prisoner's Dilemma

- GTFT = S(p, q) with

$$p = 1 \quad \text{and} \quad q = \min \left\{ 1 - \frac{T - R}{R - S}, \frac{R - P}{T - P} \right\}$$

- This is the highest level of forgiveness, q, that is still resistant against invasion by ALLD.
- Among all reactive strategies that can resist ALLD, GTFT leads to the highest payoff for the population adopting it.

Memory-one strategies

- We now consider strategies that decide between C or D based on both the opponent's and one's own last move.
- The conditional probabilities to cooperate given that the last round was CC, CD, DC, DD are p_1, p_2, p_3, p_4 .
- The Markov chain for a game between $S(p_1, p_2, p_3, p_4)$ and $S'(p'_1, p'_2, p'_3, p'_4)$ is defined by

$$\begin{array}{c}
 \begin{array}{cccc}
 & CC & CD & DC & DD \\
 \begin{array}{c} CC \\ CD \\ DC \\ DD \end{array} & \left(\begin{array}{cccc}
 p_1 p'_1 & p_1 (1 - p'_1) & (1 - p_1) p'_1 & (1 - p_1) (1 - p'_1) \\
 p_2 p'_3 & p_2 (1 - p'_3) & (1 - p_2) p'_3 & (1 - p_2) (1 - p'_3) \\
 p_3 p'_2 & p_3 (1 - p'_2) & (1 - p_3) p'_2 & (1 - p_3) (1 - p'_2) \\
 p_4 p'_4 & p_4 (1 - p'_4) & (1 - p_4) p'_4 & (1 - p_4) (1 - p'_4)
 \end{array} \right)
 \end{array}
 \end{array}$$

Old friends

- Each memory-one strategy is a point in $[0, 1]^4$.

For example:

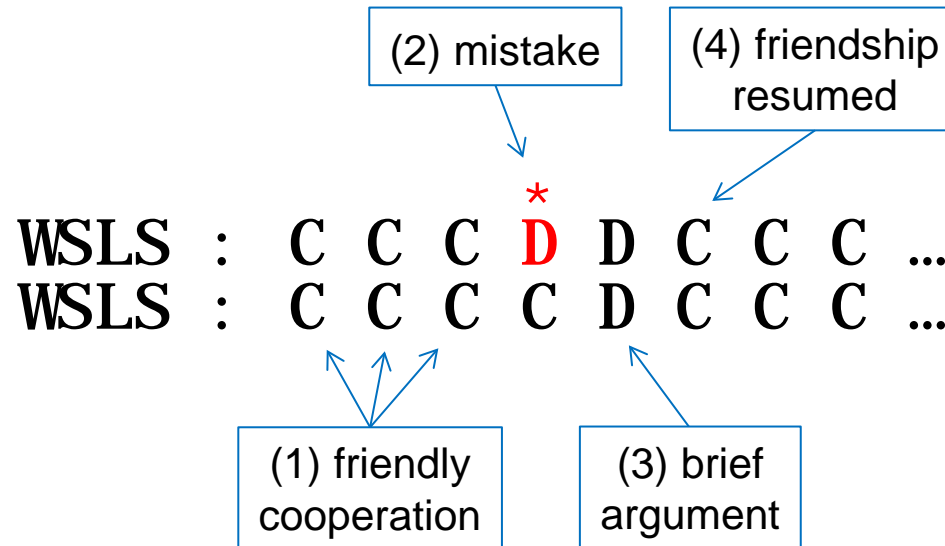
- ALLD = $S(0, 0, 0, 0)$
- ALLC = $S(1, 1, 1, 1)$
- TFT = $S(1, 0, 1, 0)$
- GTFT = $(1, 1/3, 1, 1/3)$
- Reactive strategies = $\{S(p_1, p_2, p_3, p_4) \mid p_1 = p_3, p_2 = p_4\}$

Win-stay, lose-shift (WSLS)

- $WSLS = S(1, 0, 0, 1)$, i.e., cooperate after CC or DD, defect after CD or DC.
- WSLS stays with high payoffs T or R, and shifts with low payoffs P or S.
- WSLS is stable against invasion by ALLD if $R > (T + P)/2$, but:

WSLS	:	C	D	C	D	C	D	C	...
ALLD	:	D	D	D	D	D	D	D	...

WSLS can correct mistakes



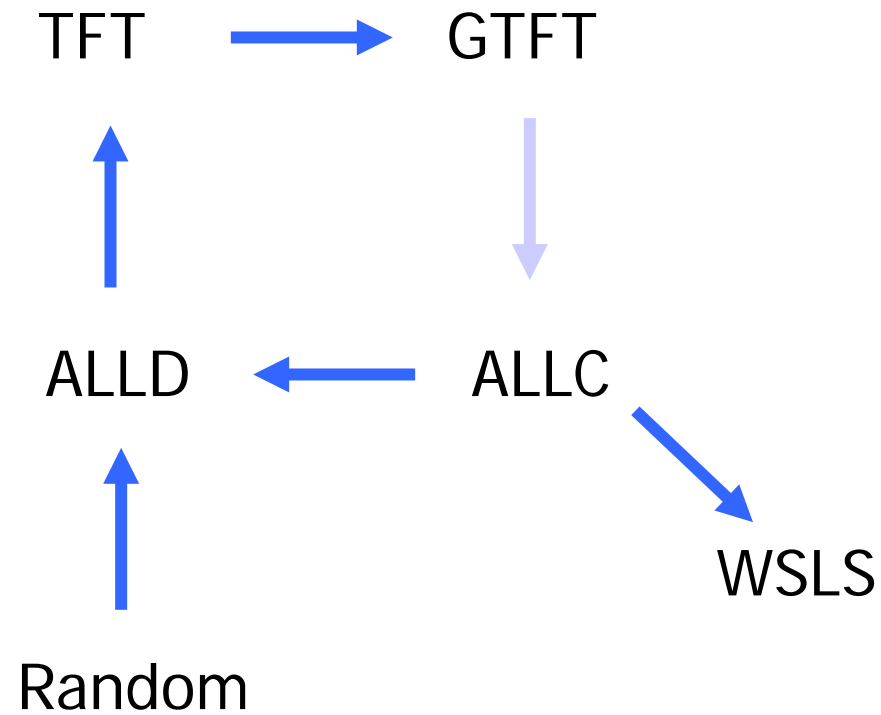
- WSLS is a deterministic corrector, whereas GTFT is a stochastic corrector.

WSLS dominates ALLC

WSLS	:	C	C	C	[*] D	D	D	D	D	...
ALLC	:	C	C	C	C	C	C	C	C	...

- “Unconditional cooperators are exploited.”
- Note that GTFT does not dominate ALLC.

War and peace



Further reading

- Nowak MA. *Evolutionary Dynamics*, Chapters 4, 5.
- Maynard Smith J. *Evolution and the Theory of Games*.
- Sinervo B, Lively CM (1996). The Rock-Paper-Scissors Game and the evolution of alternative male strategies. [Nature 380:240-243](#).