

Recap: Stability Selection

- ▶ I^* random sub-sample of size $\lfloor n/2 \rfloor$ of $\{1, \dots, n\}$
- ▶ compute $\hat{S}_\lambda(I^*)$
- ▶ repeat B times to obtain $\hat{S}_\lambda(I^{*1}), \dots, \hat{S}_\lambda(I^{*B})$
- ▶ consider the “overlap” among $\hat{S}_\lambda(I^{*1}), \dots, \hat{S}_\lambda(I^{*B})$

regarding the latter, for example:

$$\hat{\Pi}_K(\lambda) = \mathbb{P}^*[K \subseteq \hat{S}_\lambda(I^*)] \approx B^{-1} \sum_{b=1}^B I(K \subseteq \hat{S}_\lambda(I^{*b}))$$

e.g. $\hat{\Pi}_j(\lambda)$ ($j \in \{1, \dots, p\}$)

the probability \mathbb{P}^* is with respect to subsampling


The stability regularization path

Riboflavin data: $n = 115$, $p = 4088$

Y : log-production rat of riboflavin by bacillus subtilis

X : gene expressions of bacillus subtilis

all X -variables permuted except 6 “a-priori relevant” genes



left: Lasso regularization path (red: the 6 non-permuted “relevant” genes)

right: Stability path with $\hat{\Pi}_j$ on y -axis (red: the 6 non-permuted “relevant” variables stick out much more clearly from the noise covariates)

Theorem 10.1

Assume:

- ▶ exchangeability condition:
 $\{1(j \in \hat{S}(\lambda)), j \in S_0^c\}$ is exchangeable for all $\lambda \in \Lambda$
- ▶ \hat{S} is not worse than random guessing

$$\frac{\mathbb{E}[|S_0 \cap \hat{S}_\Lambda|]}{\mathbb{E}[|S_0^c \cap \hat{S}_\Lambda|]} \geq \frac{|S_0|}{|S_0^c|}.$$

Then, for $\pi_{\text{thr}} \in (1/2, 1)$:


$$\mathbb{E}[V] \leq \frac{1}{2\pi_{\text{thr}} - 1} \frac{q_\Lambda^2}{p}$$

$$q_\Lambda = \mathbb{E}[\hat{S}_\Lambda(I^*)]$$


An illustration for graphical modeling

$p = 160$ gene expressions, $n = 115$

GLasso estimator, selecting among the $\binom{p}{2} = 12'720$ features
stability selection with $\mathbb{E}[V] \leq v_0 = 30$



with permutation (empty graph is correct)



Stability Selection is extremely easy to use
and super-generic

the sufficient assumption (far from necessary) for mathematical
guarantees are restrictive
but the method seems to work well in practice

P-values based on multi sample splitting

(Ch. 11 in Bühlmann and van de Geer (2011))

Stability Selection

- ▶ uses subsampling many times – a good thing!
- ▶ provides control of the expected number of false positives rather than e.g. the familywise error rate \leadsto we will “address” this with multi sample splitting and aggregation of P-values

Fixed design linear model

$$Y = X\beta^0 + \varepsilon$$

instead of de-biased/de-sparsified method, consider the “older” technique (which is not statistically optimal but more generic and more in the spirit of stability selection)

split the sample into two parts I_1 and I_2 of equal size $\lfloor n/2 \rfloor$

- ▶ use (e.g.) Lasso to select variables based on I_1 : $\hat{S}(I_1)$
- ▶ perform low-dimensional statistical inference on I_2 based on data $(x_{I_2}^{(\hat{S}(I_1))}, Y_{I_2})$;

for example using the t -test for single coefficients β_j^0

(if $j \notin \hat{S}(I_1)$), assign the p-value 1 to the hypothesis

$H_{0,j} : \beta_j^0 = 0$,

due to independence of I_1 and I_2 , this is a “valid” strategy (see later)


validity of the (single) data splitting procedure
 consider testing $H_{0,j} : \beta_j^0 = 0$ versus $H_{A,j} : \beta_j^0 \neq 0$
 assume Gaussian errors for the fixed design linear model :
 thus, use the t -test on the second half of the sample I_2 to get a
 p-value

$$P_{\text{raw},j} \text{ from } t\text{-test based on } X_{I_2, \hat{S}(I_1)}, Y_{I_2}$$

$P_{\text{raw},j}$ is a valid p-value (controlling type I error) for testing $H_{0,j}$
 is $\hat{S}(I_1) \supseteq S_0$ (i.e., the screening property holds)
 if the screening property does not hold: $P_{\text{raw},j}$ is still valid for
 $H_{0,j}(M) : \beta_j(M) = 0$ where $M = \hat{S}(I_1)$ is a selected sub-model
 and $\beta(M) = (X_M^T X_M)^{-1} X_M^T \mathbb{E}[Y]$

a p-value lottery depending on **the random split** of the data

motif regression $n = 287$, $p = 195$



~> should aggregate/average over multiple splits!

Multiple testing and aggregation of p-values

the issue of multiple testing:

$$\tilde{P}_j = \begin{cases} P_{\text{raw},j} \text{ based on } \mathbf{Y}_{I_2}, \mathbf{X}_{I_2, \hat{S}(I_1)} & , \text{ if } j \in \hat{S}(I_1), \\ 1 & , \text{ if } j \notin \hat{S}(I_1) \end{cases}$$

thus, we can have at most $|\hat{S}(I_1)|$ false positives

\leadsto can correct with Bonferroni with factor $|\hat{S}(I_1)|$ (instead of factor p) to control the familywise error rate

$$\tilde{P}_{\text{corr},j} = \min(\tilde{P}_j \cdot |\hat{S}(I_1)|, 1) \quad (j = 1, \dots, p)$$

decision rule: reject $H_{0,j}$ if and only if $\tilde{P}_{\text{corr},j} \leq \alpha$

$\leadsto \text{FWER} \leq \alpha$

the issue with P-value aggregation:

if we run sample splitting B times, we obtain P-values

$$\tilde{p}_{\text{corr},j}^{[1]}, \dots, \tilde{p}_{\text{corr},j}^{[B]}$$

how to aggregate these dependent p-values to a single one?

for $\gamma \in (0, 1)$ define

$$Q_j(\gamma) = \min \left\{ q_\gamma(\{ \tilde{p}_{\text{corr},j}^{[b]} / \gamma; b = 1, \dots, B \}), 1 \right\},$$

where $q_\gamma(\cdot)$ is the (empirical) γ -quantile function

Proposition 11.1 (Bühlmann and van de Geer, 2011)

For any $\gamma \in (0, 1)$, $Q_j(\gamma)$ are P-values which control the FWER

example: $\gamma = 1/2$

aggregate the p-values with the sample median and multiply by the factor 2




avoid choosing γ :

$$P_j = \min \left\{ \underbrace{(1 - \log \gamma_{\min})}_{\text{price to optimize over } \gamma} \inf_{\gamma \in (\gamma_{\min}, 1)} Q_j(\gamma), 1 \right\} \quad (j = 1, \dots, p).$$




Theorem 11.1 (Bühlmann and van de Geer (2011))

For any $\gamma_{\min} \in (0, 1)$, P_j are P-values which control the FWER

$n = 100, p = 100$



$n = 100, p = 1000$



one can also adapt the method to control the False Discovery Rate (FDR)

multi sample splitting and p-value construction:

- ▶ is very generic, also for “any other” model class
- ▶ is powerful in terms of multiple testing correction: we only correct for multiplicity from $|\hat{S}(I_1)|$ variables
- ▶ it relies in theory on the screening property of the selector in practice: it is a quite competitive method!

current MSc thesis of Christoph Schultheiss: can improve multi sample splitting by multi carve methods, based on “technology” from selected inference